Lab 3 - Assignment

MM16B023[1]

[a]*Indian Institute of Technology Madras*

Keyword: VM, Cloud Functions, DataFlow, pipeline, Pcollections

Abstract: This paper presents the solutions to first assignment of the Big Data Laboratory course (CS4830) at *IIT Madras*. All the notations used are as according with the textbook Mining of massive data sets by Anand Rajaraman

Problem 1

Provide screenshot for logs of all the 3 tasks containing the required result - line count using VM, Cloud Functions and DataFlow, along with the python file for each task.

Solution:

a) Virtual Machine

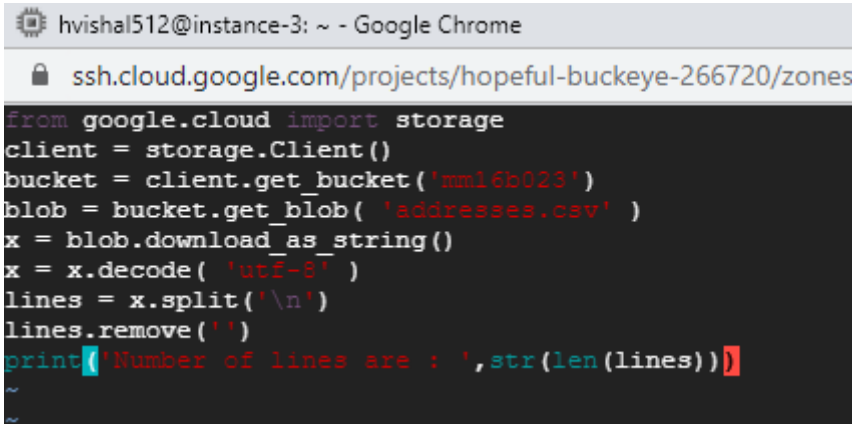

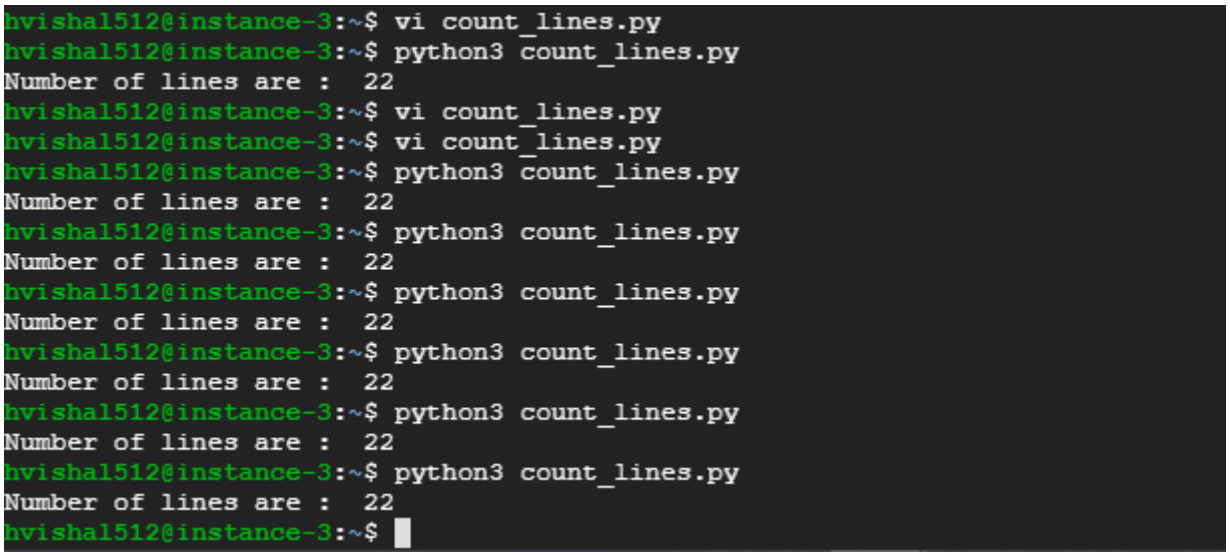

Fig1: Counting number of lines using Virtual Machine

b) Cloud Function

hvishal512@instance-3: ~/bdl3 - Google Chrome

ssh.cloud.google.com/projects/hopeful-buckeye-266720/zones

```
def cloud_func(data, context):
    from google.cloud import storage
    client = storage.Client()
    bucket = client.get_bucket('mm16b023')
    blob = bucket.get_blob(data['name'])
    x = blob.download_as_string()
    x = x.decode('utf-8')
    lines = x.split('\n')
    lines.remove('')
    print('number of lines : ' + str(len(lines)))
```

```
hvishal512@instance-3:~/bdl3$ gcloud functions deploy cloud_func --runtime python37 --trigger-resource mm16b023 --trigger-event google.storage.object.finalize
Deploying function (may take a while - up to 2 minutes)...done.
availableMemoryMb: 256
entryPoint: cloud_func
eventTrigger:
  eventType: google.storage.object.finalize
  failurePolicy: {}
  resource: projects/_/buckets/mm16b023
  service: storage.googleapis.com
ingressSettings: ALLOW_ALL
labels:
  deployment-tool: cli-gcloud
name: projects/hopeful-buckeye-266720/locations/us-central1/functions/cloud_func
runtime: python37
serviceAccountEmail: hopeful-buckeye-266720@appspot.gserviceaccount.com
sourceUploadUrl: https://storage.googleapis.com/gcf-upload-us-central1-0c9822b9-59b7-481b-b8c8-3c82cd968709/32ed1604-702a-49fe-a2d9-eb5326348b2c.zip?GoogleAccessId=service-392
399817705@gcf-admin-robot.iam.gserviceaccount.com&Expires=1581857223&Signature=bCo4zA5Uw%2F9ivEHLQVENJ1QFgDi3GQ5zN%2B7VjdDtpdr%2FtMN%2BOrQCwMn45c3W9p9WhzOZ2X1DYiz%2BWTGhSGRtDp
IVLQKadkFveHqPvr0GNp%2Bkv7hhTEWyLYrHGKdmQEnhgJjH%2FO8LNdpS2Z%2BevhirrJ41xd3zvVw2ljvlwte25sokd8i4THAH%2Bh0fSmAAl15ZtoXN4qjNwBrpAAKuyuh5Yaxh%2B6aSTZk%2BoQUwVPTDP%2FL52FRRkeDpHmU
an16dsd0lxgt7JFPHXj7LMKf0xHnW0%2FwzPivH3tV8Fem8HipvgSyx2e3h0kfKYcAhim1JMgU2rp7Clu0iHaZvoC585kX81w%3D%3D
status: ACTIVE
timeout: 60s
updateTime: '2020-02-16T12:17:36Z'
versionId: '4'
hvishal512@instance-3:~/bdl3$ gsutil cp file_path gs://mm16b023
CommandException: No URLs matched: file_path
hvishal512@instance-3:~/bdl3$ gsutil cp addresses.csv gs://mm16b023
Copying file://addresses.csv [Content-Type=text/csv]...
/ [1 files][  1.1 KiB/  1.1 KiB]
Operation completed over 1 objects/1.1 KiB.
hvishal512@instance-3:~/bdl3$ gcloud functions logs read --limit 50
LEVEL  NAME        EXECUTION_ID     TIME_UTC                 LOG
D      cloud_func  991785225927526  2020-02-16 12:19:16.412  Function execution started
I      cloud_func  991785225927526  2020-02-16 12:19:19.470  number of lines : 22
D      cloud_func  991785225927526  2020-02-16 12:19:19.473  Function execution took 3062 ms, finished with status: 'ok'
hvishal512@instance-3:~/bdl3$
```

Fig 2: Counting the number of lines using Cloud function

Using the above two methods, number of lines in the file 'addresses.csv' = 22

c) Data flow

hvishal512@instance-3: ~ - Google Chrome

ssh.cloud.google.com/projects/hopeful-buckeye-266720/zones/us-central1-a/instances/instance-3?authuser

```
import apache_beam as beam
from apache_beam.io import ReadFromText
from apache_beam.io import WriteToText
from apache_beam.options.pipeline_options import PipelineOptions
from apache_beam.options.pipeline_options import GoogleCloudOptions
from apache_beam.options.pipeline_options import StandardOptions
options = PipelineOptions()
google_cloud_options = options .view_as(GoogleCloudOptions)
google_cloud_options.project = 'hopeful-buckeye-266720' #Enter your project ID
google_cloud_options.job_name = 'bdl3'
google_cloud_options.temp_location = "gs://mm16b023/tmp" # This is to store temp
#results, format is gs: //my_bucket/tmp
options.view_as(StandardOptions).runner = 'DataflowRunner'
p = beam.Pipeline( options = options )
lines = p | 'Read' >> beam.io.ReadFromText( 'gs://iitm/files/file.txt' ) |'counting lines'
23/output_beam/' )
result = p.run()
```

Note: Some portion of the image has been cropped out for aesthetics. Please refer to the code attached.

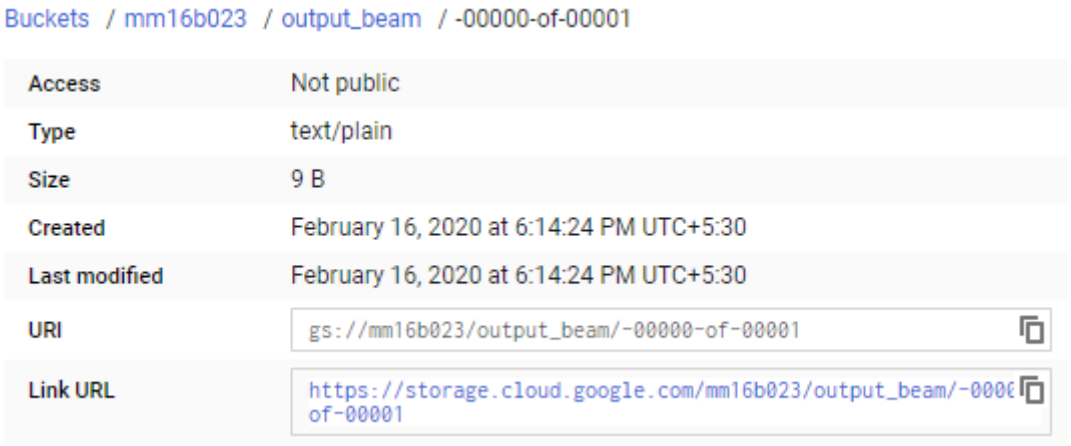

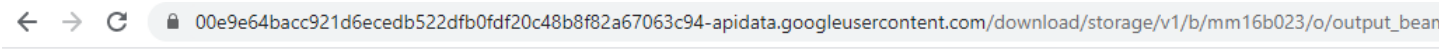


```
25974026
```

Logs

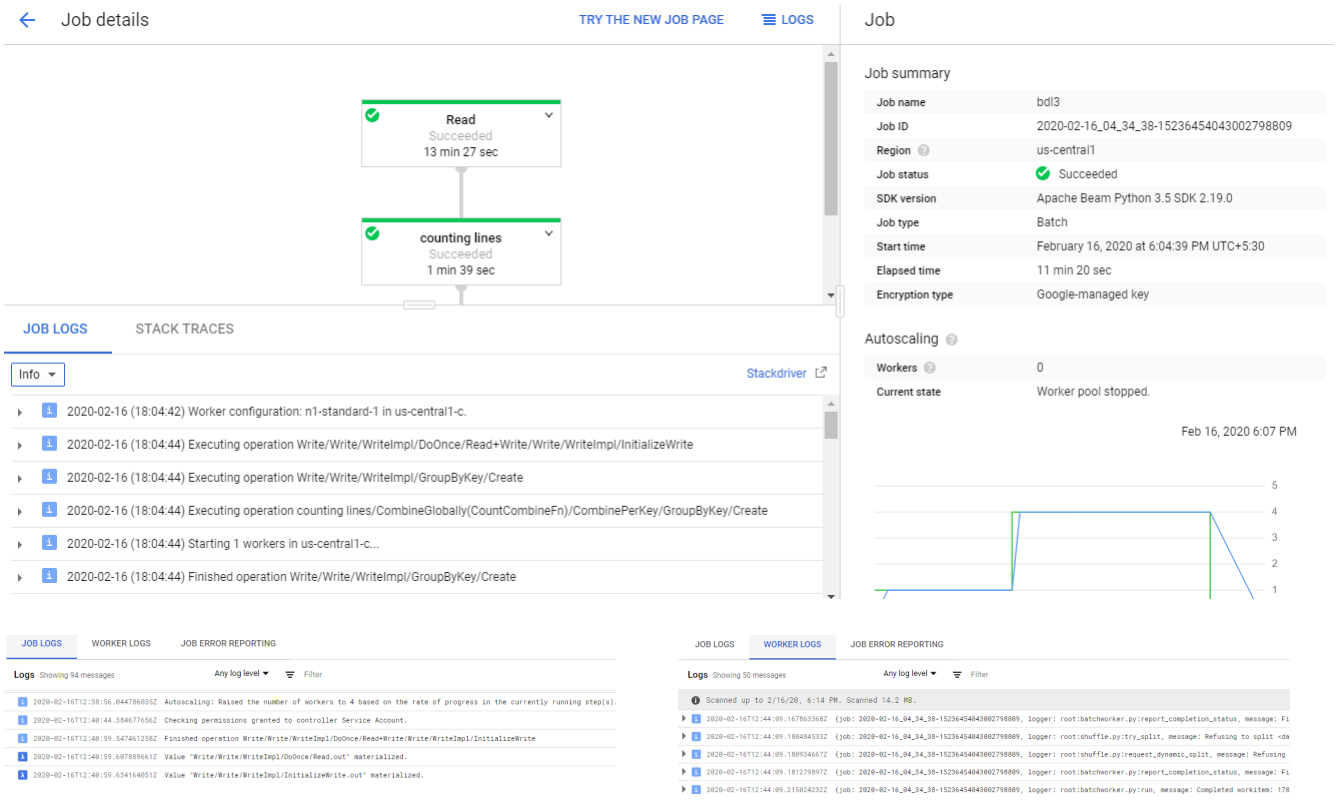

Fig3 : Counting the number of lines using DataFlow

Number of lines in the data set = 25974026

Problem2

Provide the screenshot for the execution graph created by Dataflow in the background for the pipeline object created in task 3

Solution:

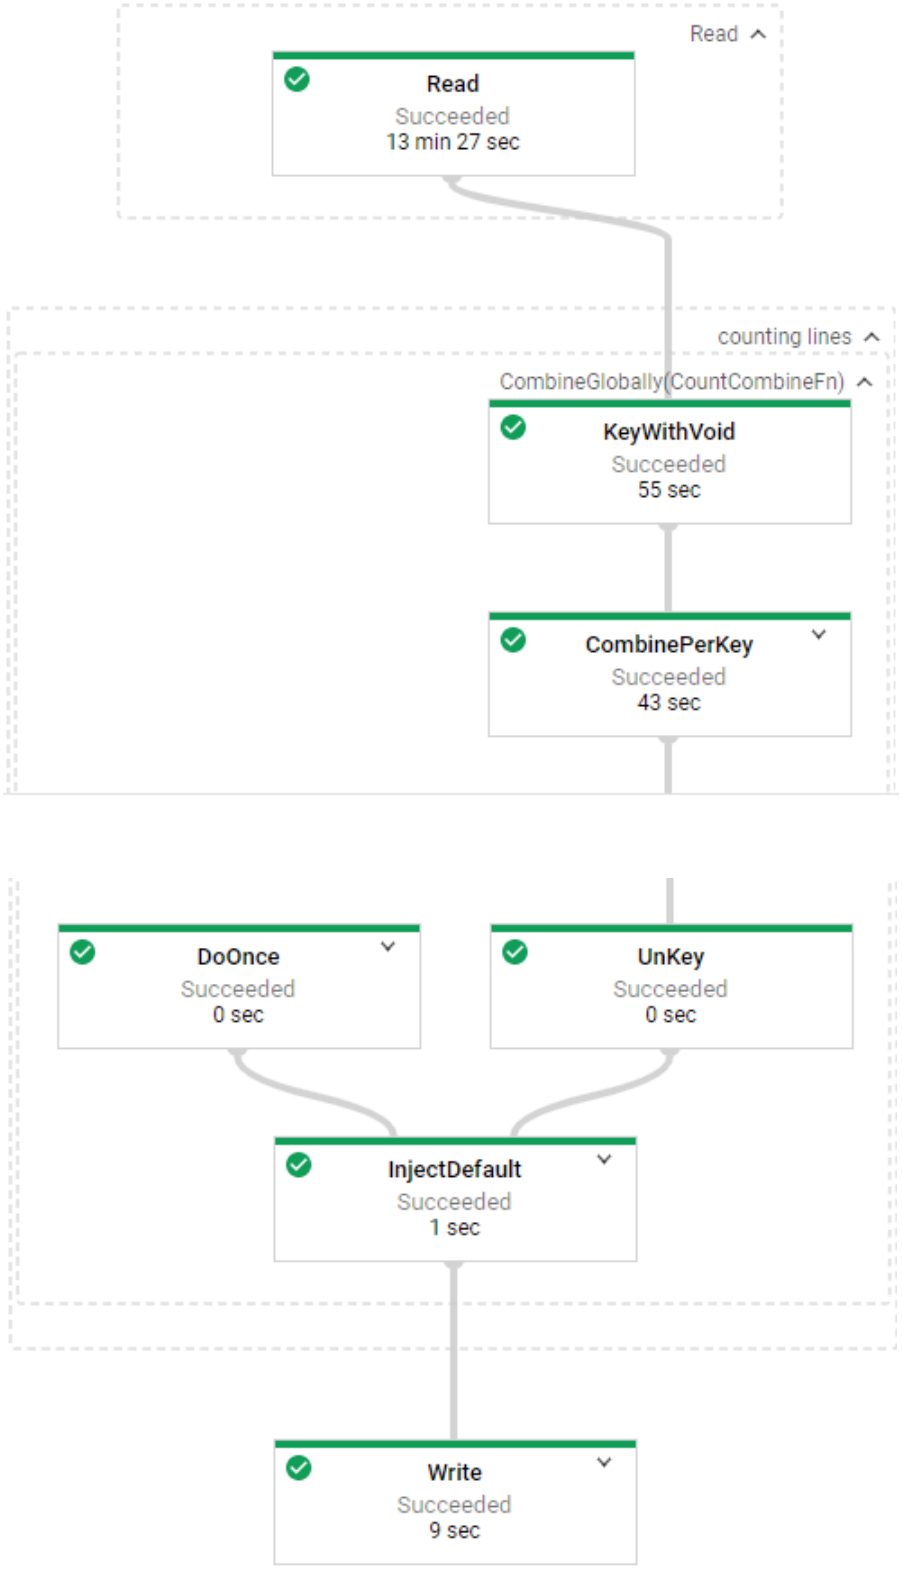

Problem 3

Explain the pipeline used in task 3. What issues did you face while trying to make the code work for task 3 and how did you resolve them?

Solution: The pipeline used has the following transforms embedded in it

- **ReadFromText** - Navigates through the file reading each line of it
- **CombineGlobally** - Stores each line as a string into a bigger dataset called PCollections
- **WriteToText** - Writes the output (no. of lines in this case) to the specified directory

Issues faced

- There was an exception thrown in the worker code when multiple lines of the code were joined by a vertical line. Surprisingly, no error came when it was written as a big single line.
- Not an issue per se, but picking the appropriate transforms and designing the right pipeline structure took some time as there are many functions and hence combinations possible.

Problem4

PCollections can handle unbounded data. What is meant by unbounded data and how do you think Pcollections can handle it? (Hint: Think on the line of triggers)

Solution:

- Data streamed from a constantly updating source is known as unbounded data. Since the data source continuously adds new elements, it has unlimited size in some sense and hence the name unbounded data.
- PCollections can handle it by modifying triggers in such a way that it maps to each collection in a streaming pipeline. This can be done using Apache Beam SDK which uses windowing (segregate unbounded data into windows) to operate the data on a combination of event time (indicated by time stamps), time of processing and the number of data elements. This modified trigger mechanism outputs whenever the watermark passes the edge of the window.
